



Large-scale hidden sweep (batch=512, seq_len=100, vanilla RNN)

A Step time vs hidden size



B Speedup over PyTorch CPU



CUDA results at batch=512 pending — run on remote GPU server:
for h in 64 128 256 512 1024 2048; do python benchmarks/run_benchmarks.py --device cuda --framework both --model vanilla --hidden \$h --seq-len 100 --batch 512; done