



Training Step Time — JAX vs PyTorch across CPU / MPS / CUDA (batch=64)

A Vanilla RNN — sequence length



B Vanilla RNN — hidden size



C Low-rank RNN — rank

